

An Abstract Application Layer
Interface to Transport Services
draft-trammell-taps-interface-01

Brian Trammell

TAPS — IETF 101 — Montreal — Tue 17 July 2018

Interface Design Principles (§3)

(a review)

We set out to define a ***single interface*** to a variety of transport protocols to be used in a variety of application design patterns, to enable applications written to a single API to make use of multiple transport protocols in terms of the features they provide, providing:

- explicit support for ***security properties*** as first-order transport features;
- ***asynchronous*** connection, transmission, and reception;
- support for ***multistreaming and multipath*** transport protocols; and
- ***atomic transmission of data***, using application-assisted framing and deframing where necessary.

Interface Diagram (as of -01)

Properties (related to Send() properties)

Require() Prefer() Ignore() Avoid() Prohibit()
Security parameters

Preconnection

Clone()

Initiate() → Ready<>
Listen() → CReceived<>
Rendezvous() → RDone<>
Stop() → Stopped<>

Endpoints

Local

Remote

Connection

Clone() → Connection Group

Send(MCtx, EOM) →
 Sent<>, Expired<>
Receive() →
 Received<Data/Metadata>
 ReceivedPartial<>

Close() → Closed<>
Abort() → Aborted<>

(non-editorial) changes since -00

- #201 Restructure Transport Properties
- #200 Rework Partial Sends and Receives
- #198 Message Receive Metadata
- #195 Ordering of API Events
- #181 Rework Interface Types
- #171 Batching Sends

(non-editorial) changes since -00

- #201 Restructure Transport Properties
- #200 Rework Partial Sends and Receives
- #198 Message Receive Metadata
- #195 Ordering of API Events
- #181 Rework Interface Types
- #171 Batching Sends

#201 Transport Parameters Rework

- All of the various ways to configure stacks (pre-selection, post-selection, and per-send) are related, but were spread throughout the document
- New approach: group all (non-security) parameters into into Properties (new §12), attempt to reclassify them.
 - Note: the authors ***do not think we have this right yet***, but we do think it's less intentionally confusing than it was.
 - Definitely needs reordering (order is kind of random)
 - May need new / renamed axes / classifications.
- Preferences still expressed using Require(), Prefer(), Avoid(), Prohibit(); send properties are bound to MessageContext passed on Send().

#201: current property "axes"

- Data type
Boolean / Enumeration / Integer / Preference
- Scope
Preconnection / Connection / Message
- Classification

| Affected Aspects | | Path & Protocol Selection | Protocol Operation | Control Flow |
|----------------------|-------------|---------------------------|--------------------|------------------|
| Level of Abstraction | Immediate | Selection Property | Protocol Property | Control Property |
| | Interpreted | Intent | | |

Properties (1/3)

| | Type | Select | protocol/ control prop. | |
|---|------|--------|----------------------------|--------------|
| | | | post- Select | per- Send |
| 12.3.1. Final | bool | | | ✓/? |
| 12.3.2. Reliable Message Transfer | pref | ✓ | | |
| 12.3.3. Configure Reliability Per Message | pref | ✓ | | |
| 12.3.4. Reliable Transfer (Message) | bool | | | ✓ |
| 12.3.5. Preservation of Data Ordering | pref | ✓ | | |
| 12.3.6. Ordered | bool | | | ✓ |
| 12.3.7. Direction of communication | enum | ? | ? | |
| 12.3.8. 0-RTT Establishment w/Idem. Send | pref | ✓ | | |
| 12.3.9. Idempotent | bool | | | ✓ |
| 12.3.10. Multistream in Group | pref | ✓ | | |
| 12.3.11. Excessive RTX Notification | pref | ✓ | | |
| 12.3.12. Exc. RTX Notification Threshold | int | ✓ | ✓ | |

Properties (2/3)

| | Type | Select | protocol/ control prop. | |
|--|-------------|--------|----------------------------|--------------|
| | | | post- Select | per- Send |
| 12.3.13. Soft Error Notification | pref | ✓ | | |
| 12.3.14. Checksum Coverage Control | pref | ✓ | | |
| 12.3.15. Checksum Coverage Length | int | | | ✓ |
| 12.3.16. Recv Checksum Requirement | int | ✓ | ✓ | |
| 12.3.17. Interface Instance / Type | (enum,pref) | ✓ | | |
| 12.3.18. PvD Instance / Type | (enum,pref) | ✓ | | |
| 12.3.19. Capacity Profile (intent) | enum | ✓ | ✓ | ✓ |
| 12.3.20. Congestion Control | pref | ✓ | | |
| 12.3.21. Niceness | int | | ✓ | ✓ |
| 12.3.22. Abort Timeout | int | ✓ | | |
| 12.3.23. Connection Group TX Scheduler | enum | ✓ | ✓ | |

Properties (3/3)

| | Type | Select | protocol/ control prop. | |
|---------------------------------------|------|--------|----------------------------|--------------|
| | | | post- Select | per- Send |
| 12.3.24. Max Idempotent Send Size | int | | r/o | |
| 12.3.25. Max No-Frag Send Size | int | | r/o | |
| 12.3.26. Max (non-partial?) Send Size | int | | r/o | |
| 12.3.27. Max (non-partial?) Recv Size | int | | r/o | |
| 12.3.28. PR Send Lifetime | int | | | ✓ |

Some Observations from the Editor (+discussion)

- Calling these axes is a little misleading: they're not orthogonal
- We have only six distinct kinds of thing:
 - Preference used for selection, scoped to preconnection, read-only after connection.
 - Property used to control how messages are sent, scoped to message (usually boolean, usually linked to selection preference).
 - Property used to control protocol operation, scoped to preconnection + connection (usually integer, controls timeout/buffer size/scheduler), possibly also usable for selection.
 - Property used to inspect protocol operation, scoped to connection, read-only (usually integer, buffer size).
 - Two enumeration/preference tuples for selecting interface/PvD.
 - One intent for path selection and scheduler configuration.

#200 Partial Send and Receive

- API is organized around atomic write/read of messages
 - (using application-supplied deframing when the underlying transport doesn't do framing, see §8.4)
- But sometimes you have a message (or a real stream) that won't fit into a buffer.
- Solution: partial read/write
 - Introduce optional EOM parameter to `Send()`; calls with `EOM = false` → still writing to a partial message identified by a given `MessageContext`.
 - `ReceivedPartial<>` event fires when a partial message is received.
- Partial read/write boundaries are not preserved.

Open issue: API for idempotent Send on establishment (#112 / #124)

- How does the application tell the stack that it wants to send some 0RTT data?
 - Some tradeoffs here, but mainly a bikeshed.
- Option 1: as in #124, hold any data sent until an explicit `Connection.Start()` call.
 - `Send()` before `Start()` is 0RTT if idempotent.
 - `Start()` is always required, even if you don't know what 0RTT is.
- Option 3: 0RTT behavior is implied by 0RTT selection properties.
 - When `Initiate()` is called and selects a 0RTT-capable stack, the actual initiation is delayed slightly to wait for the first `Send()`, which is 0RTT if idempotent.
 - Note this makes racing 0RTT-capable and 0RTT-incapable stacks impossible.
- Option 3.5: as 3, but with a `Preconnection.InitiateNow()` to override the wait-for-`Send()` behavior (e.g. for application protocols where the server sends first)
- Option 5: Add `Preconnection.Send()`, which initiates with 0RTT data.

